

Problem Statement :-

Analysis of Black Friday Sales

The analysis focuses on the sales of different products based on their Purchases, City, Gender, Age, Marital Status using *Hadoop*, *map-reduce*, *hive*, *pig*, *mongo dB*, *shell-script* and used *Excel* for graphical representation.

Dataset - <https://www.kaggle.com/mehdidag/black-friday>

Map-reduce Patterns used :-

- Secondary Sorting pattern
- Numerical Summarizations pattern
- Counting with Counters pattern
- Filtering pattern
- Top ten pattern
- Bottom ten pattern
- Distinct pattern
- Join Pattern
- Partitioning pattern

Java Map-Reduce Programs :-

1. **Pro_ProdSumPriceCityWise** - It calculates the total purchase amount of products city wise.

```
hadoop jar Pro_ProdSumPriceCityWise-0.0.1-SNAPSHOT.jar
Project/Pro_ProdSumPriceCityWise/Driver /Project/Input/BlackFriday.csv
/Project/Output/SecSort/Input/.
```

Output :-

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/SecSort/Input/* | head -20
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr
b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
P00000142      A      2700335
P00000142      B      10192027
P00000142      C      22784190
P00000242      A      36191966
P00000242      B      51435388
P00000242      C      67942452
P00000342      A      84768684
P00000342      B      102099807
P00000342      C      119868254
P00000442      A      137805778
P00000442      B      155935863
P00000442      C      174145483
P00000542      A      192547543
P00000542      B      211345432
P00000542      C      230346271
P00000642      A      251308449
P00000642      B      275397576
P00000642      C      302033993
P00000742      A      329045026
P00000742      B      356713249
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar$
```

We can see for productid P00000142 city = A collected 2700335\$ of purchase amount, city = B collected 10192027\$ of purchase amount whereas city = C collected 22784190\$

2. **Pro_SecSortProdPrice** - Secondary Sorting pattern. It uses the output file produced by the Pro_ProdSumPriceCityWise map-reduce program and provides output for each product with city having highest purchase total at top and lowest at bottom.

```
hadoop jar Pro_SecSortProdPrice-0.0.1-SNAPSHOT.jar Project/Pro_SecSortProdPrice/Driver  
/Project/Output/SecSort/Input/part-r-00000 /Project/Output/SecSort/Output/.
```

Output

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/SecSort/output/* | head -20  
WARNING: An illegal reflective access operation has occurred  
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/l  
b5.Config.getInstance()  
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.K  
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
WARNING: All illegal access operations will be denied in a future release  
prodID=P00000142, city=C, Amount=22784190|  
prodID=P00000142, city=B, Amount=10192027|  
prodID=P00000142, city=A, Amount=2700335|  
prodID=P00000242, city=C, Amount=67942452|  
prodID=P00000242, city=B, Amount=51435388|  
prodID=P00000242, city=A, Amount=36191966|  
prodID=P00000342, city=C, Amount=119868254|  
prodID=P00000342, city=B, Amount=102099807|  
prodID=P00000342, city=A, Amount=84768684|  
prodID=P00000442, city=C, Amount=174145483|  
prodID=P00000442, city=B, Amount=155935863|  
prodID=P00000442, city=A, Amount=137805778|  
prodID=P00000542, city=C, Amount=230346271|  
prodID=P00000542, city=B, Amount=211345432|  
prodID=P00000542, city=A, Amount=192547543|  
prodID=P00000642, city=C, Amount=302033993|  
prodID=P00000642, city=B, Amount=275397576|  
prodID=P00000642, city=A, Amount=251308449|  
prodID=P00000742, city=C, Amount=384785625|  
prodID=P00000742, city=B, Amount=356713249|  
cat: Unable to write to output stream.  
akash@akash-VirtualBox:~/BigData/Project_jar$
```

We can see that for ProductID = P00000142 , City = C has highest purchases whereas City = A has lowest.

3. **Pro_ProdSummary** – Numerical Summarizations pattern. It provides Summary of purchases for each product like Average, Mean, Standard deviation, Max, Min, Total and count of that product.

```
hadoop jar Pro_ProdSummary-0.0.1-SNAPSHOT.jar Project/Pro_ProdSummary/Driver  
/Project/Input/BlackFriday.csv /Project/Output/ProdSummary/.
```

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/ProdSummary/* | head -20  
WARNING: An illegal reflective access operation has occurred  
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bi  
b5.Config.getInstance()  
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.Kerberos  
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
WARNING: All illegal access operations will be denied in a future release  
0.0 0.0 NaN 0.0 0.0 0.0 1  
P00000142 11143.507079646017 10925.0 2307.94279510495 13716.0 2725.0 1.2592163E7 1130  
P00000242 10552.293800539084 10023.0 3179.231273299195 16503.0 3214.0 3914901.0 371  
P00000342 5299.928571428572 5286.0 1834.6178297875344 8896.0 1727.0 1261383.0 238  
P00000442 4795.358695652174 5194.0 1664.088124921866 8891.0 1764.0 441173.0 92  
P00000542 5419.308219178082 5375.0 1742.058891142271 8903.0 1730.0 791219.0 146  
P00000642 14913.23828125 15510.0 3633.1425203600147 19705.0 3950.0 7635578.0 512  
P00000742 6033.441176470588 5444.0 1631.819453430037 8902.0 1723.0 1435959.0 238  
P00000842 10008.722222222223 9987.0 2664.654462843017 13322.0 3347.0 360314.0 36  
P00000942 10614.5 9921.0 4371.054282176767 19154.0 4022.0 573183.0 54  
P00001042 13739.238866396761 15230.0 3346.6481344668687 19670.0 3847.0 6787184.0 494  
P00001142 6681.475524475524 6973.0 1513.412060279189 8903.0 1814.0 3821804.0 572  
P00001242 4255.828282828283 3735.0 2027.1855831619548 8841.0 1720.0 421327.0 99  
P00001342 8498.688524590163 8387.0 2762.266065757024 13693.0 2813.0 518420.0 61  
P00001442 7383.481927710844 7883.0 1617.6387968605459 9945.0 3925.0 612829.0 83  
P00001542 9243.911764705883 8146.0 4743.403239300427 19411.0 3811.0 628586.0 68  
P00001642 7570.851002865329 7908.0 1738.2658208521943 10075.0 2135.0 2642227.0 349  
P00001742 7985.144766146993 7999.0 1801.1111411668269 10081.0 2013.0 3585330.0 449  
P00001842 14420.315789473685 16180.0 4012.9336773671416 20956.0 4089.0 1095944.0 76  
P00001942 4314.588516746411 3729.0 2076.661256463911 8857.0 1713.0 901749.0 209  
cat: Unable to write to output stream.  
akash@akash-VirtualBox:~/BigData/Project_jar$
```

We can see productid P00000142 on an average was purchased at 11143.50\$ with max purchase amount of 13716\$ and min of 2725 and so on.

4. **Pro_BestCity4Product** – It provides the city with the highest purchases for all the respective products along with the purchase amount.

```
hadoop jar Pro_BestCity4Product-0.0.1-SNAPSHOT.jar Project/Pro_BestCity4Product/Driver  
/Project/Output/SecSort/Input/part-r-00000 /Project/Output/BestCity4Prod/.
```

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/BestCity4Prod/* | tail -20  
WARNING: An illegal reflective access operation has occurred  
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/u  
b5.Config.getInstance()  
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.ut  
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations  
WARNING: All illegal access operations will be denied in a future release  
P0097942      C      27941735423867  
P0098042      C      27956730094720  
P0098142      C      27971725666701  
P0098242      C      27986727653943  
P0098342      C      28001738997116  
P0098442      C      28016756229612  
P0098542      C      28031775068803  
P0098642      C      28046794802707  
P0098742      C      28061815448422  
P0098842      B      28071829575603  
P0098942      C      28086851436144  
P0099042      C      28101875303001  
P0099142      C      28116900038153  
P0099242      C      28131928515571  
P0099342      C      28146964741095  
P0099442      C      28162009560687  
P0099642      C      28177057207071  
P0099742      C      28192106784680  
P0099842      C      28207158756398  
P0099942      C      28222211717239  
akash@akash-VirtualBox:~/BigData/Project_jar$
```

For ProductId P0097942 was very popular in City C.

5. **getDetailsOfProduct.sh** – It is a shell-script that takes input ProductID from user and uses the output file generated from above map-reduce programs to give the summary details of the given Product.

```
akash@akash-VirtualBox:~/BigData/Project_jar$ cat getDetailsOfProduct.sh  
echo "Enter the product id"  
read prodID  
aaa=`hadoop fs -cat /Project/Output/ProdSummary/* | grep $prodID`  
bbb=`hadoop fs -cat /Project/Output/SecSort/Output/* | grep $prodID`  
ccc=`hadoop fs -cat /Project/Output/BestCity4Prod/* | grep $prodID`  
if [ $? -eq 0 ]  
then  
    echo "  
    echo "$(tput setaf 1)-----Summary of Product-----"  
    echo "  
    echo "$(tput setaf 1)ProductID Avg_Price          Median StdDev_in_Price    Max    Min    Total    Count"  
    echo $aaa  
    echo "  
    echo "-----City Wise Collection-----"  
    echo "  
    echo $bbb  
    echo "  
    echo "-----Best City for Sale of this Product-----"  
    echo "  
    echo $ccc | awk '{print "City : " $2 " , Amount : " $3}'  
else  
    echo "*****Product of the given ID doesnt exist as of now*****"  
fi  
akash@akash-VirtualBox:~/BigData/Project_jar$
```

Output

```
akash@akash-VirtualBox:~/BigData/Project_jar$ ./getDetailsOfProduct.sh
Enter the product id
P00069042
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/h
b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/h
b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/h
b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release

-----Summary of Product-----
ProductID Avg_Price      Median StdDev_In_Price  Max    Min    Total    Count
P00069042 11863.760180995476 13270.0 2224.4314720942975 13716.0 2648.0 2621891.0 221

-----City Wise Collection-----
prodID=P00069042, city=C, Amount=930334617518| prodID=P00069042, city=B, Amount=929378898817| prodID=P00069042, city=A, Amount=928424057376


-----Best City for Sale of this Product-----
City : C, Amount : 930334617518
akash@akash-VirtualBox:~/BigData/Project_jar$
```

We can see numerical summary as well as city wise collections and the best city for selling the product with productid P00069042.

6. **Pro_SummationCounters** – Counting with Counters pattern. It uses in build counters of map-reduce frame work to give total sales collections of each city.

```
hadoop jar Pro_SummationCounters-0.0.1-SNAPSHOT.jar Project/Pro_SummationCounters/Driver  
/Project/Input/BlackFriday.csv /Project/Output/SummationCounter/.
```

```
City
A=1295668797
B=2083431612
C=1638567969
Total=5017668378
Shuffle Errors
BAD_ID=0
CONNECTION=0
IO_ERROR=0
WRONG_LENGTH=0
WRONG_MAP=0
WRONG_REDUCE=0
File Input Format Counters
  Bytes Read=24956107
File Output Format Counters
  Bytes Written=0
A      1295668797
B      2083431612
C      1638567969
Total  5017668378
akash@akash-VirtualBox:~/BigData/Project_jar$
```



We can see city B made the highest collections during black Friday sale.

7. **Pro_PartitionerCity** – Partitioning pattern. It partitions the black Friday sales data based on the city so that we can carry out fast and detailed analysis of product city wise.

```
hadoop jar Pro_PartitionerCity-0.0.1-SNAPSHOT.jar Project/Pro_PartitionerCity/Driver
/Project/Input/BlackFriday.csv /Project/Output/CityPartitioner/.
```

City A data

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/CityPartitioner/part-r-00000 | head -20
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/hadoop/common/lib/kerberos-1.2.1.jar) method org.apache.hadoop.security.authentication.util.KerberosUtil.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
A      1001448,P00146742,F,26-35,3,A,1,0,1,13,14,15859
A      1001448,P00129042,F,26-35,3,A,1,0,1,11,,11991
A      1001448,P00000942,F,26-35,3,A,1,0,1,,7703
A      1001448,P00139242,F,26-35,3,A,1,0,8,,9786
A      1001448,P00126142,F,26-35,3,A,1,0,8,16,,9925
A      1001448,P00037242,F,26-35,3,A,1,0,8,,8089
A      1001448,P00044142,F,26-35,3,A,1,0,8,16,,6085
A      1001448,P00003642,F,26-35,3,A,1,0,8,,9762
A      1001448,P00003942,F,26-35,3,A,1,0,5,14,,7060
A      1000478,P00000142,M,51-55,16,A,2,0,3,4,5,13505
A      1000478,P00278542,M,51-55,16,A,2,0,8,,5942
A      1000477,P00217542,M,36-45,14,A,1,0,1,2,11,11665
A      1000477,P00241242,M,36-45,14,A,1,0,1,11,15,7855
A      1000477,P00218442,M,36-45,14,A,1,0,1,,11421
A      1000477,P00162442,M,36-45,14,A,1,0,1,16,,11955
A      1000477,P0098342,M,36-45,14,A,1,0,1,2,,11728
A      1000477,P00233442,M,36-45,14,A,1,0,1,2,15,12087
A      1000477,P00057642,M,36-45,14,A,1,0,1,15,16,19237
A      1000477,P00025442,M,36-45,14,A,1,0,1,2,9,19142
A      1004543,P00014542,M,26-35,2,A,0,0,8,,7847
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar$
```

```
akash@akash-VirtualBox:~/BigData/Project_jar/Project-Output$ hadoop fs -cat /Project/Output/CityPartitioner/part-r-00000 | egrep -h "B|C" | wc -l
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/hadoop/common/lib/kerberos-1.2.1.jar) method org.apache.hadoop.security.authentication.util.KerberosUtil.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
0
akash@akash-VirtualBox:~/BigData/Project_jar/Project-Output$
```

City B data

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/CityPartitioner/part-r-00001 | head -20
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/hadoop/common/lib/kerberos-1.2.1.jar) method org.apache.hadoop.security.authentication.util.KerberosUtil.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
B      1005427,P00310842,M,36-45,19,B,0,0,6,8,10,15977
B      1004734,P00345842,M,51-55,1,B,1,1,2,8,14,13082
B      1004734,P00111042,M,51-55,1,B,1,1,15,,20924
B      1004732,P00093342,M,26-35,14,B,2,1,1,8,18,11535
B      1004732,P00117242,M,26-35,14,B,2,1,8,17,,7802
B      1004732,P00085242,M,26-35,14,B,2,1,8,14,17,10036
B      1004732,P00051442,M,26-35,14,B,2,1,8,17,,9877
B      1004732,P00124642,M,26-35,14,B,2,1,16,,16610
B      1004732,P00113842,M,26-35,14,B,2,1,6,8,16,12510
B      1004732,P00041842,M,26-35,14,B,2,1,6,8,,12651
B      1004732,P00219942,M,26-35,14,B,2,1,6,16,,16486
B      1004732,P00005042,M,26-35,14,B,2,1,6,16,,19933
B      1004732,P00103042,M,26-35,14,B,2,1,1,14,,15785
B      1004732,P00255942,M,26-35,14,B,2,1,1,16,,15556
B      1004732,P00209842,M,26-35,14,B,2,1,5,8,,6978
B      1004732,P00319042,M,26-35,14,B,2,1,5,,8659
B      1004732,P00216242,M,26-35,14,B,2,1,3,4,5,10974
B      1004732,P00318742,M,26-35,14,B,2,1,1,8,,12086
B      1004732,P00184942,M,26-35,14,B,2,1,1,8,17,19217
B      1004732,P00217542,M,26-35,14,B,2,1,1,2,11,11879
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar$
```



```

akash@akash-VirtualBox:~/BigData/Project_jar/Project-OutPut$ hadoop fs -cat /Project/Output/CityPartitioner/part-r-00001 | egrep -h "A|C" | wc -l
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/hadoop/common/b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
0
akash@akash-VirtualBox:~/BigData/Project_jar/Project-OutPut$ █

```

City C data

```

akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/CityPartitioner/part-r-00002 | head -20
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/hadoop/common/b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
C 1004670,P00177242,M,36-45,5,C,2,0,12,,712
C 1004669,P00112542,M,36-45,7,C,0,0,1,11,11836
C 1004669,P00178042,M,36-45,7,C,0,0,8,14,,7825
C 1004669,P00251242,M,36-45,7,C,0,0,5,11,,6942
C 1001693,P00031142,M,26-35,12,C,1,1,16,,16162
C 1001693,P00274042,M,26-35,12,C,1,1,1,,19054
C 1001693,P00120242,M,26-35,12,C,1,1,8,,9894
C 1001693,P0097242,M,26-35,12,C,1,1,1,16,,19428
C 1001693,P00115642,M,26-35,12,C,1,1,8,14,,7981
C 1001693,P00252442,M,26-35,12,C,1,1,8,,9833
C 1001693,P00351142,M,26-35,12,C,1,1,1,8,17,11679
C 1001693,P00289842,M,26-35,12,C,1,1,1,16,,3796
C 1000214,P00257042,M,18-25,20,C,1,0,5,14,,5182
C 1000214,P00184842,M,18-25,20,C,1,0,5,6,,3673
C 1000215,P00259942,M,36-45,14,C,1,1,16,,20194
C 1000215,P00350442,M,36-45,14,C,1,1,2,3,15,15957
C 1000218,P00328442,M,36-45,14,C,3,1,5,14,,8709
C 1000218,P00331042,M,36-45,14,C,3,1,5,8,,7168
C 1000226,P00115542,M,36-45,1,C,1,0,8,,9740
C 1000229,P00184942,M,18-25,10,C,1,0,1,8,17,15786
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -lsr /Project/Output/CityPartitioner/
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/hadoop/common/b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
lsr: DEPRECATED: Please use 'ls -R' instead.
-rw-r--r-- 1 akash supergroup 0 2018-12-07 05:33 /Project/Output/CityPartitioner/_SUCCESS
-rw-r--r-- 1 akash supergroup 6846799 2018-12-07 05:33 /Project/Output/CityPartitioner/part-r-00000
-rw-r--r-- 1 akash supergroup 10740780 2018-12-07 05:33 /Project/Output/CityPartitioner/part-r-00001
-rw-r--r-- 1 akash supergroup 7905943 2018-12-07 05:33 /Project/Output/CityPartitioner/part-r-00002
akash@akash-VirtualBox:~/BigData/Project_jar$ █

```

```

akash@akash-VirtualBox:~/BigData/Project_jar/Project-OutPut$ hadoop fs -cat /Project/Output/CityPartitioner/part-r-00002 | egrep -h "A|B" | wc -l
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop-2.9.1/share/hadoop/common/b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
0
akash@akash-VirtualBox:~/BigData/Project_jar/Project-OutPut$ █

```

We can see that the black Friday sales data is now be partitioned into 3 files based on the City i.e. A, B and C.

Pig Map-Reduce Programs :-

1. **pig_FilterMaleFemale.pig** – Filtering pattern. It filters the male and female data and saves into separate files.

```
akash@akash-VirtualBox:~/BigData/Project_jar$ cat pig_FilterMaleFemale.pig
a = LOAD '/Project/Input/BlackFriday.csv' using PigStorage(',')
as (userID:chararray,prodID:chararray,gender:chararray,age:chararray,
    ocp:chararray,city:chararray,stay:chararray,married:int,catg1:chararray,
    catg2:chararray,catg3:chararray,price:long);

b = Filter a by gender=='M';
c = Filter a by gender=='F';

store b into '/Project/Output/Pig/Male';
store c into '/Project/Output/Pig/Female';
akash@akash-VirtualBox:~/BigData/Project_jar$
```

pig pig_FilterMaleFemale.pig

Output

Male data

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/Pig/Male/* | head -20
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop
b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
1000002 P00285442 M 55+ 16 C 4+ 0 8 7969
1000003 P00193542 M 26-35 15 A 3 0 1 2 15227
1000004 P00184942 M 46-50 7 B 2 1 1 8 19215
1000004 P00346142 M 46-50 7 B 2 1 1 15 15854
1000004 P0097242 M 46-50 7 B 2 1 1 16 15686
1000005 P00274942 M 26-35 20 A 1 1 8 7871
1000005 P00251242 M 26-35 20 A 1 1 5 5254
1000005 P00014542 M 26-35 20 A 1 1 8 3957
1000005 P00031342 M 26-35 20 A 1 1 8 6073
1000005 P00145042 M 26-35 20 A 1 1 2 5 15665
1000007 P00036842 M 36-45 1 B 1 1 1 14 11788
1000008 P00249542 M 26-35 12 C 4+ 1 1 5 15 19614
1000008 P00220442 M 26-35 12 C 4+ 1 5 14 8584
1000008 P00156442 M 26-35 12 C 4+ 1 8 9872
1000008 P00213742 M 26-35 12 C 4+ 1 8 9743
1000008 P00214442 M 26-35 12 C 4+ 1 8 5982
1000008 P00303442 M 26-35 12 C 4+ 1 8 14 11927
1000009 P00135742 M 26-35 17 C 0 0 6 8 16662
1000009 P00039942 M 26-35 17 C 0 0 8 5887
1000009 P00161442 M 26-35 17 C 0 0 5 14 6973
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar$
```

Female Data

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/Pig/Female/* | head -20
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/h
b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
1000001 P00069042 F 0-17 10 A 2 0 3 6 14 8370
1000001 P00248942 F 0-17 10 A 2 0 1 12 1422
1000001 P00087842 F 0-17 10 A 2 0 12 14 1057
1000001 P00085442 F 0-17 10 A 2 0 12 14 5378
1000006 P00231342 F 51-55 9 A 1 0 5 8 2079
1000006 P00190242 F 51-55 9 A 1 0 4 3 13055
1000006 P0096642 F 51-55 9 A 1 0 2 14 8851
1000006 P00058442 F 51-55 9 A 1 0 5 8 16352
1000010 P00085942 F 36-45 1 B 4+ 1 2 4 8 8886
1000010 P00118742 F 36-45 1 B 4+ 1 5 11 5875
1000010 P00297942 F 36-45 1 B 4+ 1 8 8854
1000010 P00266842 F 36-45 1 B 4+ 1 5 4 10946
1000010 P00032442 F 36-45 1 B 4+ 1 5 5152
1000010 P00105942 F 36-45 1 B 4+ 1 5 7089
1000010 P00182642 F 36-45 1 B 4+ 1 2 4 9 12909
1000010 P00186942 F 36-45 1 B 4+ 1 5 12 8770
1000010 P00155442 F 36-45 1 B 4+ 1 1 11 15 15212
1000010 P00221342 F 36-45 1 B 4+ 1 1 2 5 15705
1000010 P00087242 F 36-45 1 B 4+ 1 14 7947
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar$
```

We can see the dataset is filtered based on Gender i.e. M and F and then stored into 2 different file.

2. **pig_JoinMaleFemaleProduct.pig** – Distinct pattern and Join Pattern. It is used to get the list of products purchased by both the genders.

```
akash@akash-VirtualBox:~/BigData/Project_jar$ cat pig_JoinMaleFemaleProduct.pig
m = LOAD '/Project/Output/Pig/Male/part-m-00000'
as (userID:chararray,prodID:chararray,gender:chararray,age:chararray,
   ocp:chararray,city:chararray,stay:chararray,married:int,catg1:chararray,
   catg2:chararray,catg3:chararray,price:long);

f = LOAD '/Project/Output/Pig/Female/part-m-00000'
as (userID:chararray,prodID:chararray,gender:chararray,age:chararray,
   ocp:chararray,city:chararray,stay:chararray,married:int,catg1:chararray,
   catg2:chararray,catg3:chararray,price:long);

m1 = DISTINCT(foreach m generate prodID);
f1 = DISTINCT(foreach f generate prodID);
joinMF = JOIN m1 by prodID , f1 by prodID;

store joinMF into '/Project/Output/Pig/JoinMFProd';
akash@akash-VirtualBox:~/BigData/Project_jar$
```

pig pig_JoinMaleFemaleProduct.pig

Output

```
akash@akash-VirtualBox:~/BigData/Project_jar$ hadoop fs -cat /Project/Output/Pig/JoinMFProd/* | head -30
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr
b5.Config.getInstance())
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
P0009842      P0009842
P0009942      P0009942
P0093742      P0093742
P0093842      P0093842
P0093942      P0093942
P0094042      P0094042
P0094142      P0094142
P0094242      P0094242
P0094342      P0094342
P0094442      P0094442
P0094542      P0094542
P0094642      P0094642
P0094742      P0094742
P0094842      P0094842
P0094942      P0094942
P0095042      P0095042
P0095142      P0095142
P0095242      P0095242
P0095342      P0095342
P0095442      P0095442
P0095542      P0095542
P0095642      P0095642
P0095742      P0095742
P0095842      P0095842
P0095942      P0095942
P0096042      P0096042
P0096142      P0096142
P0096342      P0096342
P0096442      P0096442
P0096542      P0096542
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar$
```

This can be used to know which products are unisex

3. To get the total collections of the product in the descending order of their collections.

```
a = LOAD '/Project/Input/BlackFriday.csv' using PigStorage(',') as
(userID:chararray,prodID:chararray,gender:chararray,age:chararray,ocp:chararray,city:chararray,stay:
chararray,married:int,catg1:chararray,catg2:chararray,catg3:chararray,price:long);
```

```
b = foreach a generate prodID , price;
```

```
c = Group b by prodID;
```

```
d = foreach c Generate group, SUM(b.price) as totPrice;
```

```
e = order d by totPrice DESC;
```

```
store e into '/Project/Output/Pig/Top10Prodcuts/';
```

Output

```
akash@akash-VirtualBox:~/BigData/Project_jar/Project-OutPut$ hadoop fs -cat /Project/Output/Pig/Top10Prodcuts/* | head
WARNING: An illegal reflective access operation has occurred
WARNING: Illegal reflective access by org.apache.hadoop.security.authentication.util.KerberosUtil (file:/usr/local/bin/hadoop
b5.Config.getInstance()
WARNING: Please consider reporting this to the maintainers of org.apache.hadoop.security.authentication.util.KerberosUtil
WARNING: Use --illegal-access=warn to enable warnings of further illegal reflective access operations
WARNING: All illegal access operations will be denied in a future release
P00025442      27532426
P00110742      26382569
P00255842      24652442
P00184942      24060871
P00059442      23948299
P00112142      23882624
P00110942      23232538
P00237542      23096487
P00057642      22493690
P00010742      21865042
cat: Unable to write to output stream.
akash@akash-VirtualBox:~/BigData/Project_jar/Project-OutPut$
```

We can see ProductID P00025442 has made the highest sales.

Hive map-reduce

1. Top 10 products based on their total collections - Top ten pattern

create table BlackFriday (UserID String, ProductID String, Gender String, Age String, Occupation Int, City String, Stay_In_Current_City_Years String, Marital_Status Int, Product_Category_1 Int, Product_Category_2 Int, Product_Category_3 Int, Purchase Int) row format delimited fields terminated by ',';

```
hive> create table BlackFriday (UserID String, ProductID String, Gender String, Age String, Occupation Int, City String, Stay_In_Current_City_Years String, Marital_Status Int, Product_Category_1 Int, Product_Category_2 Int, Product_Category_3 Int, Purchase Int) row format delimited fields terminated by ',';
OK
Time taken: 0.005 seconds
```

LOAD Data local inpath '/home/akash/Downloads/BlackFriday.csv' overwrite into table BlackFriday;

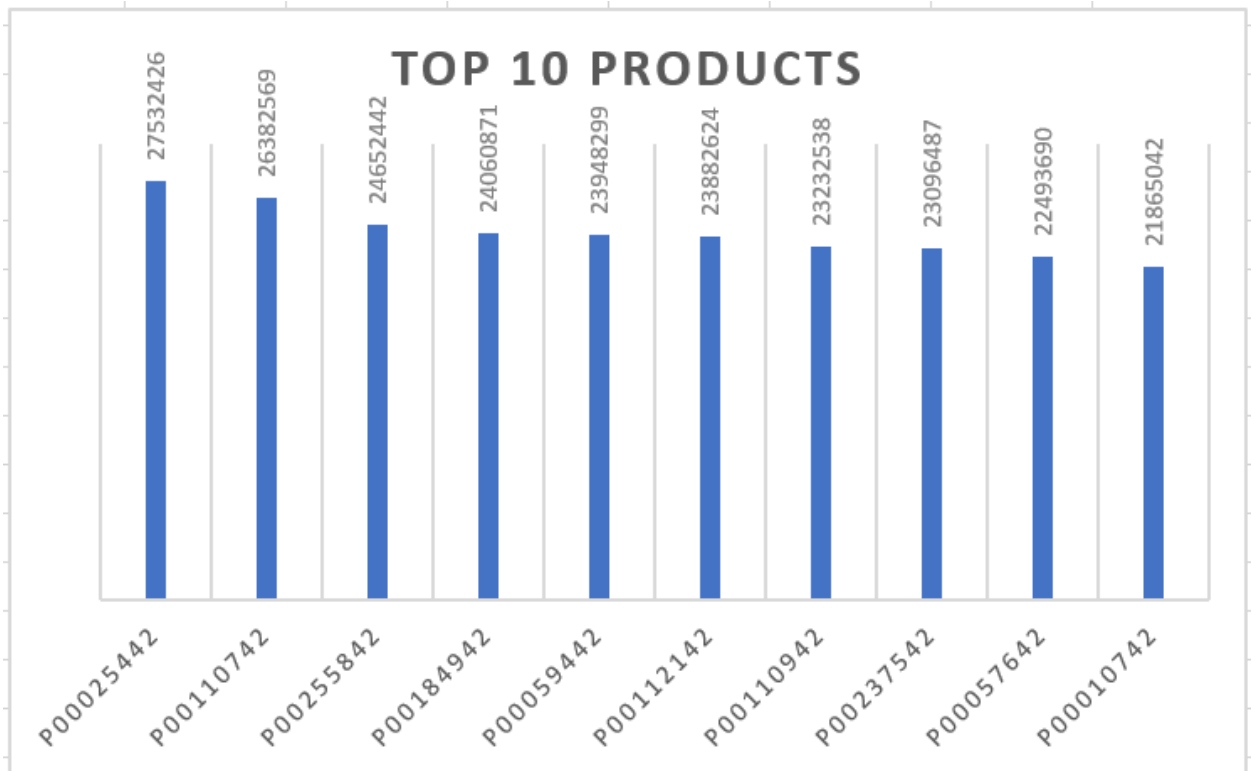
```
hive> LOAD Data local inpath '/home/akash/Downloads/BlackFriday.csv' overwrite into table BlackFriday;
Loading data to table default.blackfriday
OK
Time taken: 1.694 seconds
```

select Distinct productid, SUM(purchase) over (partition by productid ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as TotalPurchase from BlackFriday order by TotalPurchase DESC limit 10;

```
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 5.89 sec HDFS Read: 24966949 HDFS Write: 112888 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 2.93 sec HDFS Read: 117672 HDFS Write: 416 SUCCESS
Stage-Stage-3: Map: 1 Reduce: 1 Cumulative CPU: 1.96 sec HDFS Read: 5875 HDFS Write: 397 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 780 msec
OK
P00025442      27532426
P00110742      26382569
P00255842      24652442
P00184942      24060871
P00059442      23948299
P00112142      23882624
P00110942      23232538
P00237542      23096487
P00057642      22493690
P00010742      21865042
Time taken: 83.154 seconds, Fetched: 10 row(s)
hive>
```

To Store it in a file –

insert overwrite local directory '/home/akash/Downloads/Hive-out' select Distinct productid, SUM(purchase) over (partition by productid ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as TotalPurchase from BlackFriday order by TotalPurchase DESC limit 10;



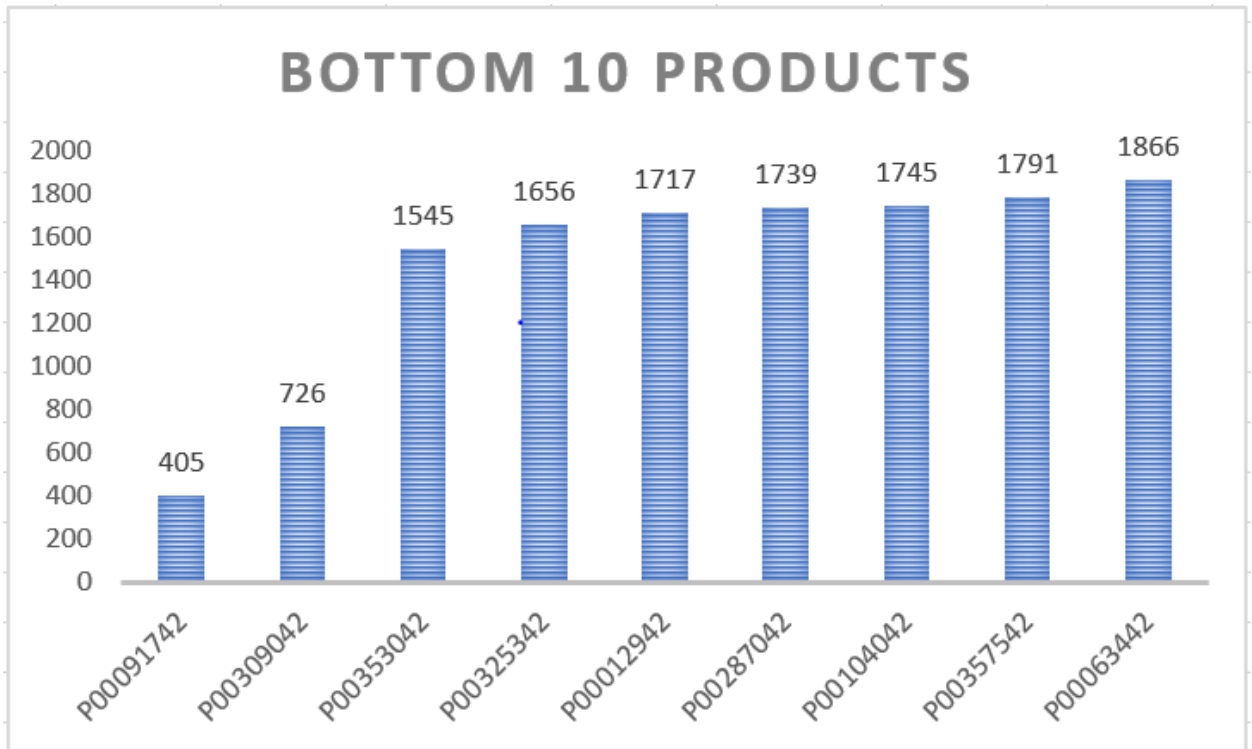
We can see ProductID P00025442 has made the highest sales.

2. Bottom 10 products based on their total collections - Bottom ten pattern

select Distinct productid, SUM(purchase) over (partition by productid ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as TotalPurchase from BlackFriday order by TotalPurchase limit 10;

```
MapReduce Jobs Launched:
Stage-Stage-1: Map: 1 Reduce: 1 Cumulative CPU: 6.36 sec HDFS Read: 24966902 HDFS Write: 112888 SUCCESS
Stage-Stage-2: Map: 1 Reduce: 1 Cumulative CPU: 2.22 sec HDFS Read: 117672 HDFS Write: 394 SUCCESS
Stage-Stage-3: Map: 1 Reduce: 1 Cumulative CPU: 2.17 sec HDFS Read: 5853 HDFS Write: 354 SUCCESS
Total MapReduce CPU Time Spent: 10 seconds 750 msec
OK
Product_ID      NULL
P00091742      405
P00309042      726
P00353042      1545
P00325342      1656
P00012942      1717
P00287042      1739
P00104042      1745
P00357542      1791
P00063442      1866
Time taken: 90.003 seconds, Fetched: 10 row(s)
hive>
```

insert overwrite local directory '/home/akash/Downloads/Hive-out/Bottom/' select Distinct productid, SUM(purchase) over (partition by productid ROWS BETWEEN UNBOUNDED PRECEDING AND UNBOUNDED FOLLOWING) as TotalPurchase from BlackFriday order by TotalPurchase limit 10;



We can see ProductID P00091742 has made the lowest sales.

MongoDB map-reduce

Dumping BlackFriday.csv into mongoDB

```
./mongoimport --db sales --collection blackfriday --type csv --headerline --file  
/c/Users/Akash/Downloads/BlackFriday.csv
```

```
Akash@DESKTOP-BBTTUHF MINGW64 /c/Program Files (x86)/mongodb-win32-x86_64-2008p1
us-ssl-3.4.17-12-g4fce55e175/bin
$ ./mongoimport --db sales --collection blackfriday --type csv --headerline --file /c/Users/Akash/Downloads/BlackFriday.csv
2018-12-13T09:35:30.100-0500    connected to: localhost
2018-12-13T09:35:32.083-0500    [#####.....] sales.blackfriday    3.32MB/23.8MB (14.0%)
2018-12-13T09:35:35.083-0500    [#####.....] sales.blackfriday    8.11MB/23.8MB (34.1%)
2018-12-13T09:35:38.084-0500    [#####.....] sales.blackfriday    13.2MB/23.8MB (55.6%)
2018-12-13T09:35:41.083-0500    [#####.....] sales.blackfriday    17.9MB/23.8MB (75.0%)
2018-12-13T09:35:43.937-0500    [#####.....] sales.blackfriday    23.8MB/23.8MB (100.0%)
2018-12-13T09:35:43.937-0500    imported 537577 documents

Akash@DESKTOP-BBTTUHF MINGW64 /c/Program Files (x86)/mongodb-win32-x86_64-2008plus-ssl-3.4.17-12-g4fce55e175/bin
```

Dumping BlackFridayA.csv, BlackFridayB.csv, BlackFridayC.csv into mongoDB i.e. the data partitioned city wise using Pro_PartitionerCity

```
Akash@DESKTOP-BBTTUHF MINGW64 /c/Program Files (x86)/mongodb-win32-x86_64-2008p1
us-ssl-3.4.17-12-g4fce55e175/bin
$ ./mongoimport --db sales --collection blackfridayA --type csv --headerline --f
ile /c/Users/Akash/Project/Project-OutPut/CityRecords.csv
2018-12-13T11:47:10.147-0500    connected to: localhost
2018-12-13T11:47:12.129-0500    [#####.....] sales.blackfridayA    3
.51MB/6.94MB (50.5%)
2018-12-13T11:47:13.853-0500    [#####.....] sales.blackfridayA    6
.94MB/6.94MB (100.0%)
2018-12-13T11:47:13.853-0500    imported 144638 documents

Akash@DESKTOP-BBTTUHF MINGW64 /c/Program Files (x86)/mongodb-win32-x86_64-2008p1
us-ssl-3.4.17-12-g4fce55e175/bin
$ ./mongoimport --db sales --collection blackfridayB --type csv --headerline --file /c/Users/Akash/Project/Project-OutPut/CityBRecords.csv
2018-12-13T11:47:37.981-0500    connected to: localhost
2018-12-13T11:47:39.969-0500    [#####.....] sales.blackfridayB    3.90MB/10.9MB (35.8%)
2018-12-13T11:47:42.970-0500    [#####.....] sales.blackfridayB    9.86MB/10.9MB (90.5%)
2018-12-13T11:47:43.473-0500    [#####.....] sales.blackfridayB    10.9MB/10.9MB (100.0%)
2018-12-13T11:47:43.473-0500    imported 226493 documents

Akash@DESKTOP-BBTTUHF MINGW64 /c/Program Files (x86)/mongodb-win32-x86_64-2008plus-ssl-3.4.17-12-g4fce55e175/bin
$ ./mongoimport --db sales --collection blackfridayC --type csv --headerline --file /c/Users/Akash/Project/Project-OutPut/CityRecords.csv
2018-12-13T11:48:01.516-0500    connected to: localhost
2018-12-13T11:48:03.509-0500    [#####.....] sales.blackfridayC    4.05MB/8.02MB (50.5%)
2018-12-13T11:48:05.634-0500    [#####.....] sales.blackfridayC    8.02MB/8.02MB (100.0%)
2018-12-13T11:48:05.634-0500    imported 166446 documents

Akash@DESKTOP-BBTTUHF MINGW64 /c/Program Files (x86)/mongodb-win32-x86_64-2008plus-ssl-3.4.17-12-g4fce55e175/bin
$
```


1. Count the number of products purchased by males and females

Over all the cities

```
> mapfunc
function () { emit(this.Gender,1); }
> redfunc
function (key, values){
var count = Array.sum(values);
return count;
}
> db.blackfriday.mapReduce(mapfunc,redfunc,{out:"ProductCount4MaleFemale"});
{
  "result" : "ProductCount4MaleFemale",
  "timeMillis" : 5961,
  "counts" : {
    "input" : 537577,
    "emit" : 537577,
    "reduce" : 10551,
    "output" : 2
  },
  "ok" : 1
}
> db.ProductCount4MaleFemale.find();
{ "_id" : "F", "value" : 132197 }
{ "_id" : "M", "value" : 405380 }
```

NUMBER OF PRODUCTS PURCHASED


F, 132197, 25%



M, 405380,
75%


For City = A

```
> db.blackfridayA.mapReduce(mapfunc,redufunc,{out:"ProductCount4MaleFemale_CityA"});
{
  "result" : "ProductCount4MaleFemale_CityA",
  "timeMillis" : 711,
  "counts" : {
    "input" : 144638,
    "emit" : 144638,
    "reduce" : 2859,
    "output" : 2
  },
  "ok" : 1
}
> db.ProductCount4MaleFemale_CityA.find()
{"_id" : "F", "value" : 34807 }
{"_id" : "M", "value" : 109831 }
```




For City B

```
> db.blackfridayB.mapReduce(mapfunc,redufunc,{out:"ProductCount4MaleFemale_CityB"});
{
  "result" : "ProductCount4MaleFemale_CityB",
  "timeMillis" : 1122,
  "counts" : {
    "input" : 226493,
    "emit" : 226493,
    "reduce" : 4370,
    "output" : 2
  },
  "ok" : 1
}
> db.ProductCount4MaleFemale_CityB.find()
{ "_id" : "F", "value" : 56494 }
{ "_id" : "M", "value" : 169999 }
```



For City C

```
db.blackfridayC.mapReduce(mapfunc,redufunc,{out:"ProductCount4MaleFemale_CityC"});
{
  "result" : "ProductCount4MaleFemale_CityC",
  "timeMillis" : 769,
  "counts" : {
    "input" : 166446,
    "emit" : 166446,
    "reduce" : 3329,
    "output" : 2
  },
  "ok" : 1
}
db.ProductCount4MaleFemale_CityC.find()
{"_id" : "F", "value" : 40896 }
{"_id" : "M", "value" : 125550 }
```




We can see that in overall as well as city wise the percentage of products brought by males and females is almost constant i.e around M-75% and F-25%

2. Total amount spent made by males and females

Overall collections

```
> mapfunc1
function () { emit(this.Gender,this.Purchase); }
> redufunc
function (key, values){
var count = Array.sum(values);
return count;
}
> db.blackfriday.mapReduce(mapfunc1,redufunc,{out:"PurchasesMadeByMaleFemale"});
{
  "result" : "PurchasesMadeByMaleFemale",
  "timeMillis" : 7307,
  "counts" : {
    "input" : 537577,
    "emit" : 537577,
    "reduce" : 10551,
    "output" : 2
  },
  "ok" : 1
}
> db.PurchasesMadeByMaleFemale.find();
{ "_id" : "F", "value" : 1164624021 }
{ "_id" : "M", "value" : 3853044357 }
```




In City A

```
> db.blackfridayA.mapReduce(mapfunc1,reducefunc,{out:"PurchasesMadeByMaleFemale_CityA"});
{
  "result" : "PurchasesMadeByMaleFemale_CityA",
  "timeMillis" : 814,
  "counts" : {
    "input" : 144638,
    "emit" : 144638,
    "reduce" : 2859,
    "output" : 2
  },
  "ok" : 1
}
> db.PurchasesMadeByMaleFemale_CityA.find()
{ "_id" : "F", "value" : 300411276 }
{ "_id" : "M", "value" : 995257521 }
```

CITY A-AMOUNT SPENT BY MALES AND FEMALES

F, 300411276, 23%




M, 995257521, 77%

In City B


```
> db.blackfridayB.mapReduce(mapfunc1,redufunc,{out:"PurchasesMadeByMaleFemale_CityB"});
{
  "result" : "PurchasesMadeByMaleFemale_CityB",
  "timeMillis" : 1256,
  "counts" : {
    "input" : 226493,
    "emit" : 226493,
    "reduce" : 4370,
    "output" : 2
  },
  "ok" : 1
}
> db.PurchasesMadeByMaleFemale_CityB.find()
{ "_id" : "F", "value" : 485312751 }
{ "_id" : "M", "value" : 1598118861 }
```

CITY B-AMOUNT SPENT BY MALES AND FEMALES



In City C

```
> db.blackFridayC.mapReduce(mapfunc1,redufunc,{out:"PurchasesMadeByMaleFemale_CityC"});
{
  "result" : "PurchasesMadeByMaleFemale_CityC",
  "timeMillis" : 941,
  "counts" : {
    "input" : 166446,
    "emit" : 166446,
    "reduce" : 3329,
    "output" : 2
  },
  "ok" : 1
}
> db.PurchasesMadeByMaleFemale_CityC.find()
{"_id" : "F", "value" : 378899994 }
{"_id" : "M", "value" : 1259667975 }
```




We can see Male spent more as compared to females in all the cities.

3. Total amount spent by different age groups

Overall


```
> mapfunc2
function () { emit(this.Age, this.Purchase); }
> db.blackfriday.mapReduce(mapfunc2, redfunc, {out: "AmountSpentByDiffAgeGroups"});
{
  "result" : "AmountSpentByDiffAgeGroups",
  "timeMillis" : 3196,
  "counts" : {
    "input" : 537577,
    "emit" : 537577,
    "reduce" : 26862,
    "output" : 7
  },
  "ok" : 1
}

> db.AmountSpentByDiffAgeGroups.find()
{ "_id" : "0-17", "value" : 132659006 }
{ "_id" : "18-25", "value" : 901669280 }
{ "_id" : "26-35", "value" : 1999749106 }
{ "_id" : "36-45", "value" : 1010649565 }
{ "_id" : "46-50", "value" : 413418223 }
{ "_id" : "51-55", "value" : 361908356 }
{ "_id" : "55+", "value" : 197614842 }
```




In City A

```
db.blackfridayA.mapReduce(mapfunc2,redufunc,{out:"AmountSpentByDiffAgeGroups_CityA"});
{
  "result" : "AmountSpentByDiffAgeGroups_CityA",
  "timeMillis" : 859,
  "counts" : {
    "input" : 144638,
    "emit" : 144638,
    "reduce" : 6191,
    "output" : 7
  },
  "ok" : 1
}
db.AmountSpentByDiffAgeGroups_CityA.find()
{ "_id" : "0-17", "value" : 21657219 }
{ "_id" : "18-25", "value" : 240169988 }
{ "_id" : "26-35", "value" : 647728672 }
{ "_id" : "36-45", "value" : 236366595 }
{ "_id" : "46-50", "value" : 62621435 }
{ "_id" : "51-55", "value" : 57155885 }
{ "_id" : "55+", "value" : 29969003 }
```




In City B

```
> db.blackfridayB.mapReduce(mapfunc2,redufunc,{out:"AmountSpentByDiffAgeGroups_CityB"});
{
  "result" : "AmountSpentByDiffAgeGroups_CityB",
  "timeMillis" : 1343,
  "counts" : {
    "input" : 226493,
    "emit" : 226493,
    "reduce" : 9769,
    "output" : 7
  },
  "ok" : 1
}
> db.AmountSpentByDiffAgeGroups_CityB.find()
{ "_id" : "0-17", "value" : 47513782 }
{ "_id" : "18-25", "value" : 385183031 }
{ "_id" : "26-35", "value" : 825779950 }
{ "_id" : "36-45", "value" : 426450764 }
{ "_id" : "46-50", "value" : 185012084 }
{ "_id" : "51-55", "value" : 163783897 }
{ "_id" : "55+", "value" : 49708104 }
```



In City C


```
> db.blackFridayC.mapReduce(mapfunc2,redufunc,{out:"AmountSpentByDiffAgeGroups_CityC"});
{
  "result" : "AmountSpentByDiffAgeGroups_CityC",
  "timeMillis" : 971,
  "counts" : {
    "input" : 166446,
    "emit" : 166446,
    "reduce" : 10538,
    "output" : 7
  },
  "ok" : 1
}
> db.AmountSpentByDiffAgeGroups_CityC.find()
{ "_id" : "0-17", "value" : 63488005 }
{ "_id" : "18-25", "value" : 276316261 }
{ "_id" : "26-35", "value" : 526240484 }
{ "_id" : "36-45", "value" : 347832206 }
{ "_id" : "46-50", "value" : 165784704 }
{ "_id" : "51-55", "value" : 140968574 }
{ "_id" : "55+", "value" : 117937735 }
```



We can see that most of the amount came from age group 26-35 and with highest contribution from City A for this age group i.e 50%.

4. Total Amount spread by married and unmarried customers

```
> mapfunc3
function () { emit(this.Marital_Status, this.Purchase); }
> db.blackfriday.mapReduce(mapfunc3, redfunc, { out: "AmountSpentByMaritalStatus" });
{
  "result" : "AmountSpentByMaritalStatus",
  "timeMillis" : 2826,
  "counts" : {
    "input" : 537577,
    "emit" : 537577,
    "reduce" : 10718,
    "output" : 2
  },
  "ok" : 1
}
> db.AmountSpentByMaritalStatus.find()
{ "_id" : 0, "value" : 2966289500 }
{ "_id" : 1, "value" : 2051378878 }
```



City Wise

```
> mapfunc4
function () {emit({city:this.City_Category,marital_stat:this.Marital_Status},this.Purchase);}
> db.blackfriday.mapReduce(mapfunc4,redfunc,{out:"AmountSpentByMaritalStatusCitywise"});
{
  "result" : "AmountSpentByMaritalStatusCitywise",
  "timeMillis" : 4204,
  "counts" : {
    "input" : 537577,
    "emit" : 537577,
    "reduce" : 27166,
    "output" : 6
  },
  "ok" : 1
}
> db.AmountSpentByMaritalStatusCitywise.find()
{ "_id" : { "city" : "A", "marital_stat" : 0 }, "value" : 806390391 }
{ "_id" : { "city" : "A", "marital_stat" : 1 }, "value" : 489278406 }
{ "_id" : { "city" : "B", "marital_stat" : 0 }, "value" : 1233719389 }
{ "_id" : { "city" : "B", "marital_stat" : 1 }, "value" : 849712223 }
{ "_id" : { "city" : "C", "marital_stat" : 0 }, "value" : 926179720 }
{ "_id" : { "city" : "C", "marital_stat" : 1 }, "value" : 712388249 }
```

We can see that unmarried people from city B have spent the most whereas married people of City A have spent the least of all

Conclusion

- City B made the highest sales collection followed by City C and City A.
- Males and Unmarried people spent more compared to Females and married people respectively.
- People belonging to 55+ Age group spent the least whereas people of age group 26-35 took more interest in buying stuff.
- Product-Id P00025442 made the highest sales collection whereas Product-Id P00091742 made the lowest.